Lecture 16

$\hat{\hat{\theta}}_{1\cdot} = \frac{n_{1\cdot}}{n}$ proportion of people with black hair.

HB

$\Rightarrow \hat{\hat{\theta}}_{i\cdot} = \frac{n_{i\cdot}}{n}$, $\hat{\hat{\theta}}_{\cdot j} = \frac{n_{\cdot j}}{n}$

$$\Rightarrow \hat{\phi} = \sum_{i=1}^{r} \sum_{j=1}^{c} \frac{(\theta_{ij} - n\hat{\theta}_{i\cdot}\hat{\theta}_{\cdot j})}{\underbrace{n\hat{\theta}_{i\cdot}\hat{\theta}_{\cdot j}}_{E_{ij}}} \xrightarrow{d} \chi^2_{(r-1)(c-1)}$$

a fact that could be proved in a certain time permitted

"Chi-Squared test of independence"

$\hat{\phi} = 41.20$ $\alpha = 5\%$. $F_{\chi^2_9} = (16.99) = 95\%$

$\Rightarrow$ Reject H_0

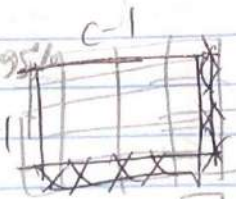

$RET = [0, 16.99]$

We conclude that hair color and eye color are dependent.

This class has focused mainly on three goals of inference: estimation, testing and confidence sets. We will continue to study these three goals but... we will also study some tangential "meta" concept that are classic.

Here is one such classic "meta concept". usually you are

(a) given a dataset $x_1, \ldots x_n$ then you

(b) assume a DGP, then you

(c) define one or (many) inferential target parameters theta, the

(d) Compute $\hat{\theta} = w(x_1 \ldots x_n)$ and then you

(e) make a CI / run a test at size α.

Let examine (b). How do you just "assume a DGP"? Sometimes you really do know the DGP e.g a coin flipped repeatedly is iid Bernoulli (θ), a die rolled repeatedly is iid uniform discrete. But what about "daily wind speeds at JFK airport" or "rat survival times" (like on midterm) or "daily percentage returns of the S&P 500"? The DGPs for the last three are very complicated and they're unknown.

What if we wanted to "guess" the DGP's? This is actually a really big part of what

Statisticians do. This is called "model fitting". DGP = model. Model you kinda make up and hopefully they are useful for whatever you are doing. Why don't we proceed as follows:

let's guess M candidate models/ DGPs $m = 1, 2, \ldots, M$ and then

(1) pick the "best" model out of my M guesses and maybe

(2) provide a weighting score to each of the M guesses (low scores indicate bad guesses and high scores indicate good guesses). Goal (1) is famous and called "Model Selection". In 342, we do a

little of this atheoretically. Here we'll do it more theoretically.

Model selection is more fundamental than you realize. It's actually the entire problem of all of science. For example, let's say you have some astronomical data on movement of different planets, stars, etc. You want to fit a model (guess a DGP) for the force on celestial bodies with masses m_1 and m_2 at a distance r from eachother (i.e. "gravity"). Consider the following models.

Model 1: $F = G\frac{m_1 m_2}{r^2}$ Newton's Law

Model 2: $F = G_1\frac{m_1 m_2}{r^2} + G_2\frac{m_1 m_2}{r^2}$ Newton's extension

Model 3: $F = G_1\frac{m_1 m_2}{r^2}e^{-G_2 r}$ Laplace extension

Which model is best? We know all these are wrong because Einstein came and disproved them with general relativity.

Lets talk about Model Selection techniques. Our data $x_1 \dots x_n$ comes from an unknown DGP. Here are M candidate models:

Model 1: $f_1(x_1 \dots x_n ; \theta_{11}, \theta_{1K}) = f(\theta_{11}, \dots, \theta_{1K_1} | x_1 \dots x_n)$

model 2: $f_2(x_1 \dots x_n ; \theta_{21} \dots \theta_{2K_2}) = f(\theta_{21} \dots, \theta_{2K_2} | x_1 \dots x_n)$

⋮

Model M_i: $f_m(x_1 \dots x_n ; \theta_{m1} \dots \theta_{mK_m}) = f(\theta_{m1} \dots \theta_{mK_m} | x_1 \dots x_n)$

K_1 is the # of Parameters in model 1,

K_2 is the num of parameter in model 2, K_m is the num of parameter into model M. Each K_m could be different.

Why don't we just select the model that has the highest likelihood?

$$m_i := \underset{m}{\operatorname{argmax}} \left\{ l_M \left(\theta_{M_1} \ldots \theta_{M_{K_m}} ; x_1 \ldots x_n \right) \right\}$$

$$= \underset{m}{\operatorname{argmax}} \left\{ l_m \left(\theta_{M_1} \ldots \theta_{M K_m} ; x_1 \ldots x_n \right) \right\}$$

The problem with this is we don't know the values of θ for any of the models!

So let's do recharadze $K_1 + K_2 + K_m$ times! We'll estimate each of the parameters

using MLE's and plug them all in and then

$$m_{\alpha} = \arg\max \left\{ \ell\left(\hat{\theta}^{MLE}_{M_1}, \ldots \hat{\theta}^{MLE}_{M k_m}, X_1 \ldots X_n \right) \right\}$$

You could do this. But... it will not give you the best model. Why?

$\ell(\hat{\theta}^{MLE}_{m_1}, \ldots \hat{\theta}^{MLE}_{m k_m}; X_1 \ldots X_n)$ is an estimate for $\ell(\theta_{m_1}, \ldots \theta_{m k_m}; X_1 \ldots X_n)$

and it's biased.... with many assumptions, you can prove that

$$\text{Bias}\left[\ell(\hat{\theta}^{MLE}_{m}, \ldots \hat{\theta}^{MLE}_{m k_m}; X_1 \ldots X_n)\right] = K_m > 0$$

There is positive Bias (meaning the log-likelihood would appear

higher on average) and this bias increases you use more parameters in your candidate models. This was figured out by H. Akaike, a Japanese statistician and he published it in 1974. Once you have the bias, you just use it to correct your estimate:

$$\ell(\theta_{m_1}, \ldots \theta_{mk_m}) \approx \ell(\hat{\theta}^{MLE}_{m_1}, \ldots \hat{\theta}^{MLE}_{mk_m}, X_1 \ldots X_n) - k_m$$

Recall that log-likelihood is always negative for discrete DUP's and almost always negative for continuous DUP's. So let's flip it's sign and multiply by 2:

Complexity penalty.

$$AIC_m = -2\ell(\hat{\theta}^{MLE}_{m_1}, \ldots \hat{\theta}^{MLE}_{mk_m}, X_1 \ldots X_n) + 2k_m$$

↑ (Akaike

(AKaike's Information Criterion)

The "best" log likelihood is the largest i.e closest to zero

So once negated, the "best" negative log likelihood is the smallest i.e closest to zero.

$m_d = \text{argmax}($